

Anjali Jain

anjali9@illinois.edu

Team - Team232

Team Name - Mocking Bird

Team Members - Single member Project

Final Project Phase-I (Summer 2024)

CS513: Theory & Practice of Data Cleaning

| |
|--|
| Overview |
| <p>For my Summer Final Project, the NYPL-menus dataset was assigned to me. This dataset consists of four tables: Dish, Menu, MenuItem, and MenuPages. It provides a detailed historical overview of dining, featuring 17,545 menus from 233 venues across 3,714 locations, with 423,397 unique dishes. Notable trends include frequent appearances of coffee, tea, and celery.</p> <p>Predominantly featuring daily menus and special occasions like anniversaries, the dataset offers insights into culinary trends, pricing patterns, and dining habits over time. Analysis revealed the top occasions, status distribution, popular dishes, and the impressive longevity of celery, which appeared from year 1 to 2928. This comprehensive collection is a valuable resource for food historians and researchers.</p> |
| About Dataset |
| <p>NYPL-menus</p> <p>a. Content: "What's on the menu?": A mix of simple bibliographic description of the menus (created by The New York Public Library) and the culinary and economic content of the menus themselves (transcribed by you).</p> <p>b. Source: http://menus.nypl.org/data</p> |
| Dataset in details |

The NYPL-menu dataset offers a fascinating glimpse into the history of dining, capturing detailed information about menus and dishes served at various events and locations over multiple years.

As mentioned in overview this dataset includes 17,545 menus from 233 venues across 3,714 locations, featuring an impressive 423,397 unique dishes. On average, each menu showcases around 76 dishes, contributing to a total of 1,332,343 dish appearances.

As we delve into the data, notable trends emerge. Coffee, tea, and celery stand out as the most frequently listed items. This rich trove of information includes metadata about the physical characteristics of the menus, the events they were created for, and the establishments that hosted these events.

The collection predominantly features daily menus, but also includes special occasions like anniversaries. This dataset has potential to provide invaluable insights into culinary trends, pricing patterns, and dining habits over time, making it a robust resource for food historians and researchers in various fields.

Using Python in a Jupyter Notebook, a detailed analysis of the Menu and Dish tables was conducted by me. The results revealed a total of 17,545 menus, representing 233 unique venues and 3,714 unique locations, with an average of 75.62 dishes per menu. Among the top occasions, daily menus were the most common, followed by complimentary/testimonial events and anniversaries.

Interestingly, the dataset showed a sparse language distribution but a clear status distribution, with 99.01% of the data marked as complete and 0.99% under review. The analysis identified 423,397 unique dishes, which appeared a total of 1,332,343 times on menus. The date range of the menus extends from year 0 to year 2928, highlighting the longevity of certain dishes.

Among the top five most popular dishes, coffee topped the list with 8,484 appearances, followed by tea, celery, olives, and radishes. Notably, celery holds the distinction of being the longest-running dish, appearing consistently from year 1 to year 2928.

In summary, the NYPL-menu dataset is a treasure trove of historical culinary data, offering deep insights into the evolution of food trends, pricing, and dining practices over an expansive timeframe.

Snapshot of from my Jupyter notebook

Run analysis & plot graph for table DISH -

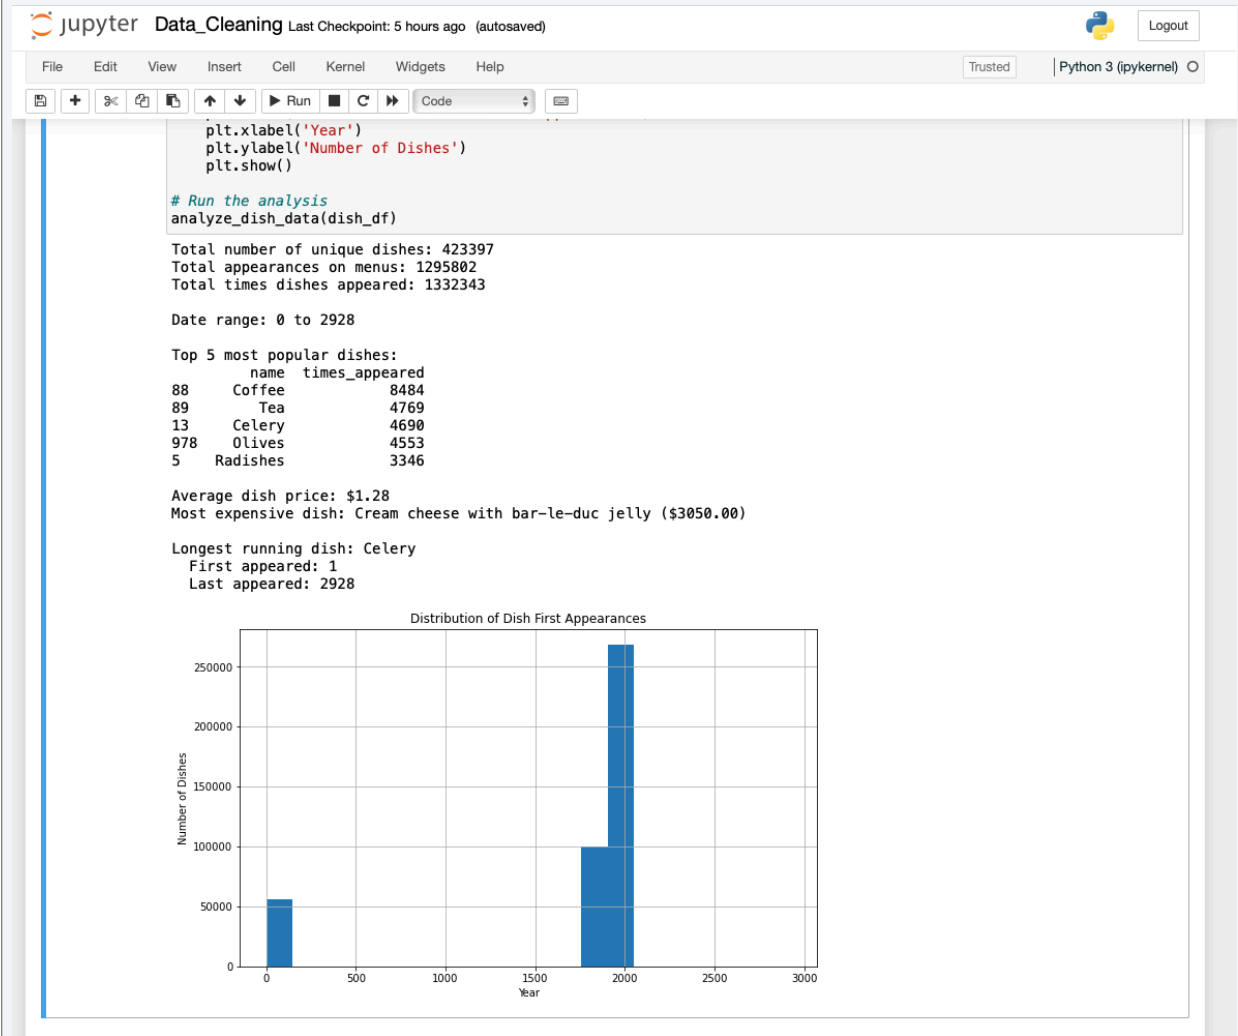


Table Fields & description

After performing initial analysis on the four CSV files—Menu, MenuPage, MenuItem, and Dish—I created four tables using Microsoft Excel. Here are the details of each table and its attributes.

| | Table Attributes | | |
|------------|------------------|---|--|
| Table Name | Field Name | Field Description | |
| Dish | ID | Unique identifier for each dish | |
| | Name | Name of the dish | |
| | Description | Description of the dish. | |
| | Menus_appeared | Number of menus on which dish has appeared | |
| | times_appeared | Total number of times the dish has appeared on all menus. | |
| | first_appeared | Year when the dish first appeared on a menu. | |
| | last_appeared | Year when the dish last appeared on a menu. | |
| | lowest_price | Lowest recorded price of the dish | |
| | highest_price | Highest recorded price of the dish | |
| | | | |

| | Table Attributes | |
|-------------------|----------------------|---|
| Table Name | Field Name | Field Description |
| Menu | ID | Unique identifier for each menu |
| | name | Name of the event or establishment |
| | Sponsor | Sponsor of the menu |
| | Event | Type of event for example is it breakfast or Dinner ? |
| | Venue | Venue of the event |
| | Place | Place (city ,state)where the event took place . |
| | physical_description | Physical description of the menu (e.g., card, booklet). |
| | Occasion | Special occasion associated with the menu. |
| | Notes | Additional notes about the menu. |
| | call_number | Call number for the menu. |
| | Keywords | Keywords associated with the menu. |
| | Language | Language of the menu. |
| | Date | Date of the event. |
| | Location | Specific location of the venue. |

| | Table Attributes | |
|--|------------------|--|
| | | the venue. |
| | location_type | Type of location (e.g., hotel, restaurant). |
| | Currency | Currency used in the menu. |
| | Currency_symbol | Symbol of the currency.(A currency symbol or currency sign is a graphic symbol used to denote a currency unit.) |
| | Status | Status of the menu (e.g., complete, under review). |
| | Page_count | Number of pages in the menu. |
| | Dish_count | Number of dishes in the menu. |
| | | |

| | Table Attributes | | |
|----------|------------------|--|--|
| Table | Field Name | Field Description | |
| Menuitem | ID | Unique identifier for each menu item | |
| | Menu_page_id | Foreign key referencing the menu page. | |
| | Price | Price of the dish. | |
| | high_price | High price of the dish. | |
| | dish_id | Foreign key referencing the dish. | |
| | created_at | Timestamp when the price was recorded. | |
| | updated_at | Timestamp when the price was last updated. | |
| | xpos | X-coordinate position of the price on the menu page. | |
| | ypos | Y-coordinate position of the price on the menu page. | |
| | | | |

| | Table Attributes | | |
|-----------------|------------------|--------------------------------------|--|
| | | | |
| Table | Field Name | Field Description | |
| MenuPage | ID | Unique identifier for each menu page | |
| | menu_id | Foreign key referencing the menu. | |
| | page_number | Page number in the menu | |
| | image_id | Image ID of the scanned menu page. | |
| | full_height | Full height of the scanned image. | |
| | full_width | Full width of the scanned image. | |
| | uuid | Unique identifier for each image | |

Database Schema SQL

Below schema defines the structure of the dataset and outlines the relationships between the tables. The Dish table captures information about individual dishes, the Menu table captures information about specific menus and events, the MenuPage table captures metadata about the scanned pages of each menu, and the MenuItem table captures the pricing information for each dish on each menu page.

Each table has primary key ID, Table MenuPage has foreign key menu_id and MenuItem has 2 foreign key dish_id and menu_page_id.

Dish -

```
CREATE TABLE Dish (  
  id INT PRIMARY KEY,  
  name VARCHAR(255),  
  description TEXT,  
  menus_appeared INT,  
  times_appeared INT,  
  first_appeared YEAR,  
  last_appeared YEAR,  
  lowest_price DECIMAL(5,2),  
  highest_price DECIMAL(5,2)  
);
```

Menu -

```
CREATE TABLE Menu (  
  id INT PRIMARY KEY,  
  name VARCHAR(255),  
  sponsor VARCHAR(255),  
  event VARCHAR(255),  
  venue VARCHAR(255),  
  place VARCHAR(255),  
  physical_description TEXT,  
  occasion VARCHAR(255),  
  notes TEXT,  
  call_number VARCHAR(255),  
  keywords TEXT,  
  language VARCHAR(255),  
  date DATE,  
  location VARCHAR(255),  
  location_type VARCHAR(255),  
  currency VARCHAR(50),  
  currency_symbol VARCHAR(5),  
  status VARCHAR(50),  
  page_count INT,  
  dish_count INT  
);
```

Menu Page -

```
CREATE TABLE MenuPage (  
  id INT PRIMARY KEY,  
  menu_id INT,  
  page_number INT,  
  image_id INT,  
  full_height INT,  
  full_width INT,  
  uuid VARCHAR(255),  
  FOREIGN KEY (menu_id) REFERENCES Menu(id)  
);
```

MenuItem -

```
CREATE TABLE MenuItem (  
  id INT PRIMARY KEY,  
  menu_page_id INT,  
  price DECIMAL(5,2),  
  high_price DECIMAL(5,2),  
  dish_id INT,  
  created_at TIMESTAMP,  
  updated_at TIMESTAMP,  
  xpos FLOAT,  
  ypos FLOAT,  
  FOREIGN KEY (menu_page_id) REFERENCES MenuPage(id),  
  FOREIGN KEY (dish_id) REFERENCES Dish(id)  
);
```

ER diagram

Below ER diagram is created from the above schemas illustrates the entities involved in a NYPL-menus, their attributes, and the relationships that exist between them.

ER Overview :

Menu->MenuPage(1,Many)
MenuPage-MenuItem(1,Many)
Dish -> MenuItem (1,Many)

Below is the ERD diagram created using LucidChart online tool.



For Historical Pricing Trend

Perform Data Cleaning

Remove any duplicate records.
To make sure currency formats are consistent (e.g., all prices in USD).
Fill in or remove missing values, particularly for prices and dates.
Correct inconsistent date formats.
Normalize dish names where slight variations exist (e.g., "Chicken Gumbo" and "Chicken gumbo" should be treated as the same dish).

Target (Main) use case U1 Diagram: data cleaning is necessary and sufficient

Obtain Clean Data

Perform Data Analysis

Group data by year and calculate average prices for each dish.
Identify price trends over the years for popular dishes.
Analyze price variations across different locations and events.

Obtain Pricing Insights

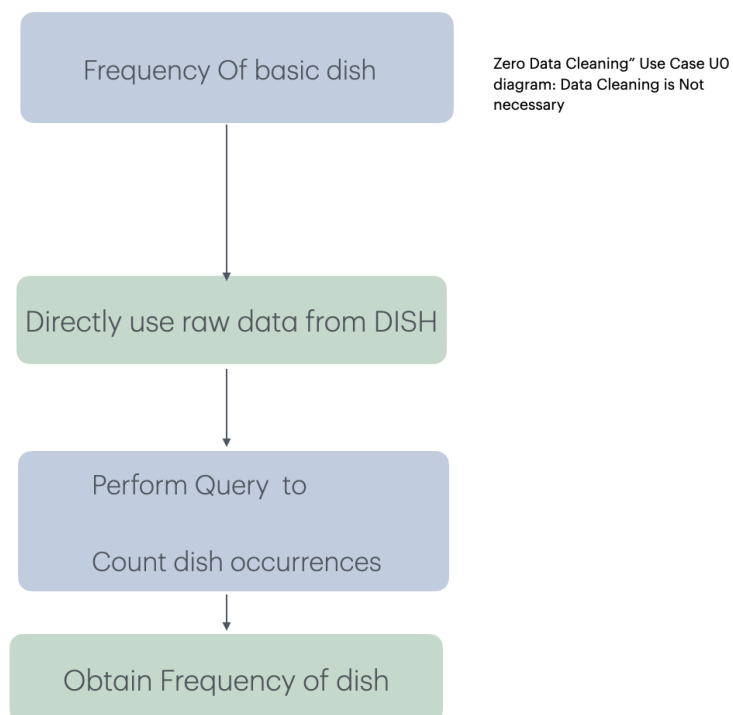
Result - Cleaned data will allow intended user for accurate and meaningful analysis of historical pricing trends, providing insights into how food prices have evolved over time.

"Zero data cleaning" use case U0: data cleaning is not necessary

Use Case U0 - To count the number of times each dish appears across all menus using the raw data

In order to count the number of times each dish appears across all menus. This use case does not require any data cleaning since the raw data is sufficient for counting occurrences.

To count the number of times each dish appears across all menus.

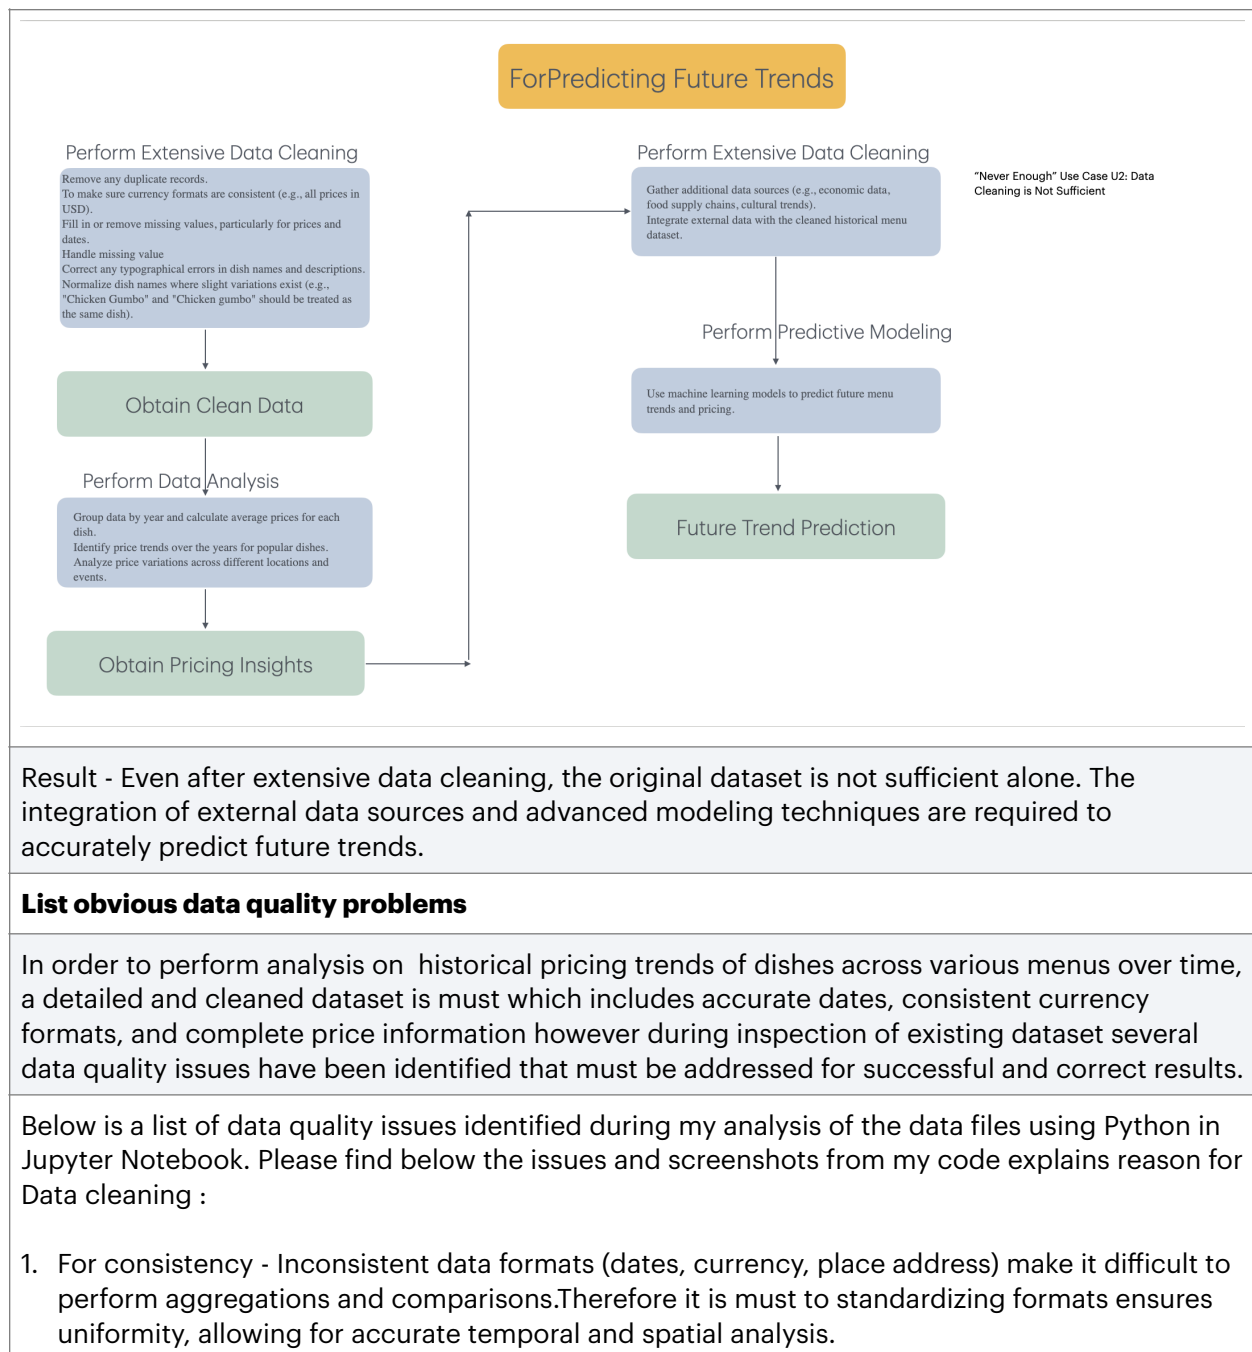


Result - The raw data of DISH is good enough to provide a count of dish occurrences without any cleaning.

"Never enough" use case U2 : data cleaning is not sufficient

Use Case U2 - Predicting Future Menu Trends and Pricing

In order to achieve prediction of future trends in menu offerings and pricing based on historical data, more comprehensive data and additional external factors are required. This use case involves complex modeling and external data sources such as economic indicators, food supply data, and cultural trends.



Jupyter Data_Cleaning Last Checkpoint: 6 hours ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3 (ipykernel)

```

import numpy as np

# First, let's examine the structure of our dataframes
def show_dataframe_info(df, name):
    print(f"\n{name} Dataframe Info:")
    print(df.info())
    print("\nSample data:")
    print(df.head())
    print("\nColumn names:")
    print(df.columns.tolist())

# Run this for each of your dataframes
# show_dataframe_info(menu_df, "Menu")
# show_dataframe_info(dish_df, "Dish")
# show_dataframe_info(item_df, "Menu Item")
# show_dataframe_info(page_df, "Menu Page")

# Now let's create functions based on the actual structure of your dataframes

# 1. Date issues (assuming 'date' is in menu_df)
def show_date_issues(df):
    if 'date' in df.columns:
        return df[pd.to_datetime(df['date'], errors='coerce').isnull()].head()
    else:
        return "No 'date' column found in this dataframe"

```

In [17]: show_date_issues(menu_df)

Out[17]:

| occasion | notes | call_number | keywords | language | date | location | location_type | currency | currency_symbol | status | page |
|---------------------------|--|-------------|----------|----------|------------|--------------------------------------|---------------|----------|-----------------|----------|------|
| NaN | REGULAR DINNER, SOUP, ONE MEAT DISH, TWO VEGET... | 1900-2517 | NaN | NaN | NaN | Mills Hotel restaurant | NaN | Cents | c | complete | |
| BREAKFAST; | MENU HANDWRITTEN; | 1900-2328 | NaN | NaN | 0190-03-06 | Quebec Steamship Company Ss Trinidad | NaN | NaN | NaN | complete | |
| NaN | NaN | NaN | NaN | NaN | NaN | Mills Hotel restaurant | NaN | Dollars | \$ | complete | |
| NaN | NaN | NaN | NaN | NaN | NaN | Russell House | NaN | NaN | NaN | complete | |
| COMPLIMENTARY/TESTIMONIAL | MENU IN FRENCH; WINES; LIST OF TOASTS & MUSIC P... | 1906-783 | NaN | NaN | NaN | Lord Mayor Of Cardiff | NaN | NaN | NaN | complete | |

2. As seen above field values are missing "NaN", inconsistent data and typographical errors can lead to incorrect calculations and trends. Therefore filling missing values, removing duplicates, and correcting errors ensures that the analysis is based on accurate data." and write for this point as reason of data cleaning Different currencies for a same dish.

3. To standardize currency values across datasets, particularly when different currencies are used to price the same dish. Without normalization, analysis may be skewed, leading to inaccurate comparisons and trends. By converting all currency values to a single standard, such as USD or EUR, data integrity is maintained, ensuring consistent and reliable analysis."

Different currencies for a same dish

```
In [19]:
def show_currency_issues_for_same_dish(menu_df, dish_df, menu_item_df, menu_page_df):
    # First, let's merge the necessary dataframes
    merged_df = pd.merge(menu_item_df, menu_page_df[['id', 'menu_id']], left_on='menu_page_id', right_on='id', suffixes=('_item', '_page'))
    merged_df = pd.merge(merged_df, dish_df[['id', 'currency', 'currency_symbol']], left_on='menu_id', right_on='id', suffixes=('_menu', '_dish'))
    merged_df = pd.merge(merged_df, dish_df[['id', 'name']], left_on='dish_id', right_on='id', suffixes=('', '_dish'))

    # Group by dish name and count unique currencies
    dish_currency_counts = merged_df.groupby('name')['currency'].nunique().reset_index()

    # Filter dishes that appear with multiple currencies
    multi_currency_dishes = dish_currency_counts[dish_currency_counts['currency'] > 1]

    # Get examples for these dishes
    examples = []
    for dish in multi_currency_dishes['name'].head():
        dish_examples = merged_df[merged_df['name'] == dish][['name', 'currency', 'currency_symbol', 'price', 'high_price']]
        examples.append(dish_examples)

    # Combine all examples
    result = pd.concat(examples).reset_index(drop=True)

    return result

# Usage
currency_issues = show_currency_issues_for_same_dish(menu_df, dish_df, item_df, page_df)
```

In [20]: currency_issues

Out[20]:

| | name | currency | currency_symbol | price | high_price |
|---|---------------------------------|------------------|-----------------|-------|------------|
| 0 | " American Rye | Belgian Francs | BEF | 0.50 | 7.5 |
| 1 | " American Rye | Francs | FF | 0.50 | 7.5 |
| 2 | "Red Brand" sirloin steak | Canadian Dollars | C\$ | 1.25 | NaN |
| 3 | "Red Brand" sirloin steak | Dollars | \$ | 1.25 | NaN |
| 4 | "Red Brand" small sirloin steak | Canadian Dollars | C\$ | 1.00 | NaN |
| 5 | "Red Brand" small sirloin steak | Dollars | \$ | 1.00 | NaN |
| 6 | "Salad" Lettuce & Tomatoes | Belgian Francs | BEF | NaN | NaN |
| 7 | "Salad" Lettuce & Tomatoes | Francs | FF | NaN | NaN |
| 8 | (1) Lambchops a la Murillo | Cents | c | 0.00 | NaN |
| 9 | (1) Lambchops a la Murillo | Dollars | \$ | 0.40 | NaN |

4. Inconsistent dish names across datasets can create ambiguity and hinder accurate analysis. Data cleaning involves standardizing dish names by resolving spelling variations, abbreviations, and synonyms. This ensures that all references to the same dish are uniform, enabling precise aggregation and analysis of culinary trends and preferences.

I took below example of chicken salad.

Out[22]:

| | name | id | menus_appeared |
|--------|---------------------|--------|----------------|
| 201 | Chicken salad | 217 | 1809 |
| 11018 | Chicken Salad. | 13723 | 8 |
| 26298 | Chicken [salad] | 33341 | 9 |
| 26699 | Chicken Salad, | 33844 | 1 |
| 31238 | Chicken (SALAD) | 40006 | 6 |
| 155858 | Chicken Salad 3 25 | 195543 | 1 |
| 286578 | Chicken Salad 1.75 | 361195 | 1 |
| 298440 | Chicken Salad | 375982 | 408 |
| 298656 | Chicken [Salad] | 376218 | 2 |
| 298915 | chicken salad | 376510 | 26 |
| 302318 | Chicken Salad | 380239 | 13 |
| 305282 | Chicken (salad) | 383517 | 8 |
| 317433 | Chicken salad | 397020 | 1 |
| 320011 | Chicken [Salad] | 399860 | 1 |
| 354394 | CHICKEN SALAD | 439070 | 8 |
| 369827 | chicken salad | 455981 | 1 |
| 378079 | chicken Salad | 464458 | 1 |
| 397279 | Chicken, Salad | 486569 | 1 |
| 409519 | Chicken [Salad] (1) | 500251 | 1 |

5. Identifying and correcting multiple outlier details is crucial in data cleaning to ensure statistical robustness and accuracy in analysis. Outliers, when left unaddressed, can skew results and misrepresent trends. Cleaning involves methods such as removing or adjusting outliers based on statistical measures like z-scores or interquartile range, thereby improving the reliability of insights derived from the data.

In [24]: multiple_outlier_details

Out[24]:

| | Item ID | Dish Name | Price | Currency | Outlier Type | Menu ID | Menu Date | Menu Venue | Page Number |
|---|---------|--------------------------|----------|--------------|--------------|---------|-----------|------------|-------------|
| 0 | 26464 | Dom Perignon | 180000.0 | Italian Lire | High | 26464 | NaN | HOTEL | 5.0 |
| 1 | 26464 | Cristal | 180000.0 | Italian Lire | High | 26464 | NaN | HOTEL | 5.0 |
| 2 | 26464 | Krug | 160000.0 | Italian Lire | High | 26464 | NaN | HOTEL | 5.0 |
| 3 | 26464 | Veuve Clicquot Ponsardin | 110000.0 | Italian Lire | High | 26464 | NaN | HOTEL | 5.0 |
| 4 | 26464 | Moet et Chandon | 100000.0 | Italian Lire | High | 26464 | NaN | HOTEL | 5.0 |

Some additional reason are -

1. For accurate outcomes - As seen above in all tables(Menu,Dish,MenuPage and MenuItem) of dataset values are missing, inconsistent data and typographical errors can lead to incorrect calculations and trends. Therefore filling missing values, removing duplicates, and correcting errors ensures that the analysis is based on accurate data.
2. Reliability - Inconsistent or incorrect data can lead to unreliable insights, affecting decision-making. Therefore Data cleaning improves the reliability of the dataset, making it fit for generating trustworthy insights.
3. Completeness - Missing descriptions or partial records can result in incomplete analysis Therefore one must Ensure all necessary data is present and complete allows for comprehensive analysis.

In conclusion, data cleaning transforms the raw dataset into a cleaned version that is accurate, consistent, reliable, and complete. This cleaned dataset is crucial for performing meaningful and accurate historical pricing trends analysis, as described in the main use case U1.

Devise an initial plan

My plan for Phase 2 is to ensure early submission to avoid any last minute hassle.

| Step Number | Activity | Responsibility | Timeline |
|-------------|--|----------------|------------|
| S1 | Description of dataset D and matching use case U1 | Anjali | 07/10/2024 |
| S2 | Profiling of D to identify the quality problems P that need to be addressed to support U1; | Anjali | 07/13/2024 |

| | | | |
|----|---|--------|------------|
| S3 | Performing the data cleaning process using one or more tools to address the problems P (here you should describe which tools you are planning to use, e.g., OpenRefine; Python; etc.) S | Anjali | 07/20/2024 |
| S4 | Checking that your new dataset D' is an improved version of D, e.g., by documenting that certain problems P are now absent and that U1 is now supported | Anjali | 07/21/2024 |
| S5 | Documenting the types and amount of changes that have been executed on D to obtain D' . | Anjali | 07/24/2024 |